

Data Science and Decision Making Assignment 1

#Brigthon Data Visualization

##Exploratory Data Analysis

```
#Brigthon Data Visualization
#Important Libraries to be used in the code
import warnings
warnings.filterwarnings('ignore')
import os
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
from matplotlib.lines import Line2D

# For seasonal decomposition
from statsmodels.tsa.seasonal import seasonal_decompose, STL
# Import additional functions needed
from statsmodels.tsa.stattools import adfuller, kpss
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.stats.diagnostic import acorr_ljungbox

#imports to get data from the GDrive
from google.colab import drive
drive.mount('/content/gdrive')

Mounted at /content/gdrive

#Mounting the path for the data
GOOGLE_DRIVE_PATH_AFTER_MYDRIVE = os.path.join('./Project
documentation-20240129/weatherdata_for_students')
GOOGLE_DRIVE_PATH = os.path.join('gdrive',
'MyDrive',GOOGLE_DRIVE_PATH_AFTER_MYDRIVE)
print('List files: ', os.listdir(GOOGLE_DRIVE_PATH))

List files: ['colchester_001.csv', 'colchester_002.csv',
'colchester_003.csv', 'colchester_004.csv', 'colchester_005.csv',
'colchester_007.csv', 'colchester_006.csv', 'colchester_008.csv',
'colchester_009.csv', 'colchester_010.csv', 'colchester_012.csv',
'colchester_011.csv', 'colchester_013.csv', 'colchester_014.csv',
'colchester_016.csv', 'colchester_015.csv', 'colchester_018.csv',
'colchester_017.csv', 'colchester_020.csv', 'colchester_019.csv',
'colchester_021.csv', 'colchester_022.csv', 'colchester_023.csv',
'colchester_024.csv', 'colchester_025.csv', 'colchester_026.csv',
```

[illegible]

'colchester_174.csv', 'colchester_176.csv', 'colchester_175.csv',
'colchester_177.csv', 'colchester_178.csv', 'colchester_179.csv',
'colchester_180.csv', 'colchester_181.csv', 'colchester_183.csv',
'colchester_184.csv', 'colchester_182.csv', 'colchester_186.csv',
'colchester_185.csv', 'colchester_187.csv', 'colchester_188.csv',
'colchester_189.csv', 'colchester_190.csv', 'colchester_191.csv',
'colchester_192.csv', 'colchester_193.csv', 'colchester_195.csv',
'colchester_196.csv', 'colchester_194.csv', 'brighton_003.csv',
'brighton_001.csv', 'brighton_002.csv', 'brighton_004.csv',
'brighton_005.csv', 'brighton_006.csv', 'brighton_009.csv',
'brighton_011.csv', 'brighton_014.csv', 'brighton_010.csv',
'brighton_013.csv', 'brighton_008.csv', 'brighton_007.csv',
'brighton_012.csv', 'brighton_015.csv', 'brighton_017.csv',
'brighton_018.csv', 'brighton_016.csv', 'brighton_019.csv',
'brighton_020.csv', 'brighton_021.csv', 'brighton_022.csv',
'brighton_024.csv', 'brighton_023.csv', 'brighton_025.csv',
'brighton_026.csv', 'brighton_027.csv', 'brighton_028.csv',
'brighton_029.csv', 'brighton_030.csv', 'brighton_031.csv',
'brighton_032.csv', 'brighton_033.csv', 'brighton_035.csv',
'brighton_034.csv', 'brighton_036.csv', 'brighton_037.csv',
'brighton_038.csv', 'brighton_039.csv', 'brighton_042.csv',
'brighton_040.csv', 'brighton_041.csv', 'brighton_043.csv',
'brighton_044.csv', 'brighton_045.csv', 'brighton_047.csv',
'brighton_046.csv', 'brighton_048.csv', 'brighton_049.csv',
'brighton_052.csv', 'brighton_055.csv', 'brighton_062.csv',
'brighton_056.csv', 'brighton_063.csv', 'brighton_061.csv',
'brighton_059.csv', 'brighton_060.csv', 'brighton_050.csv',
'brighton_051.csv', 'brighton_064.csv', 'brighton_053.csv',
'brighton_057.csv', 'brighton_058.csv', 'brighton_054.csv',
'brighton_071.csv', 'brighton_074.csv', 'brighton_067.csv',
'brighton_069.csv', 'brighton_081.csv', 'brighton_076.csv',
'brighton_068.csv', 'brighton_077.csv', 'brighton_073.csv',
'brighton_080.csv', 'brighton_066.csv', 'brighton_079.csv',
'brighton_075.csv', 'brighton_070.csv', 'brighton_072.csv',
'brighton_065.csv', 'brighton_078.csv', 'brighton_096.csv',
'brighton_097.csv', 'brighton_093.csv', 'brighton_082.csv',
'brighton_094.csv', 'brighton_083.csv', 'brighton_088.csv',
'brighton_100.csv', 'brighton_092.csv', 'brighton_091.csv',
'brighton_087.csv', 'brighton_099.csv', 'brighton_101.csv',
'brighton_089.csv', 'brighton_090.csv', 'brighton_084.csv',
'brighton_095.csv', 'brighton_098.csv', 'brighton_086.csv',
'brighton_085.csv', 'brighton_116.csv', 'brighton_119.csv',
'brighton_109.csv', 'brighton_111.csv', 'brighton_112.csv',
'brighton_118.csv', 'brighton_103.csv', 'brighton_110.csv',
'brighton_105.csv', 'brighton_117.csv', 'brighton_108.csv',
'brighton_114.csv', 'brighton_113.csv', 'brighton_107.csv',
'brighton_106.csv', 'brighton_102.csv', 'brighton_104.csv',
'brighton_115.csv', 'brighton_124.csv', 'brighton_128.csv',
'brighton_125.csv', 'brighton_127.csv', 'brighton_120.csv',

```

'brighton_123.csv', 'brighton_122.csv', 'brighton_121.csv',
'brighton_126.csv']

import os
import pandas as pd

# Specify the folder directory where the dataset is located.
df_path = GOOGLE_DRIVE_PATH

#list of all files in the path
file_list = [file for file in os.listdir(GOOGLE_DRIVE_PATH) if
file.endswith('.csv')]

brighton_df = []

# Iterate over each CSV file
for file_ in file_list:
    # Construct the full path to the CSV file
    file_path = os.path.join(df_path, file_)

    # Check for file has an index column named '0'
    index_column = pd.read_csv(file_path, nrows=1).columns[0] == '0'

    # Change header value based on the value of the column
    header = 1 if index_column else "infer"

    current_dataframe = pd.read_csv(file_path, header=header)

    # If the file name contains "Brighton", add its DataFrame to the
    list
    if file_.startswith('brighton'):
        brighton_df.append(current_dataframe)

# Concatenate all the datasets into one
f_brighton_df = pd.concat(brighton_df, ignore_index=True)

#information about the DataFrame
print(f_brighton_df.info())

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 122844 entries, 0 to 122843
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   datetime              122844 non-null object
1   temp                  122590 non-null float64
2   dew                   122568 non-null float64
3   humidity              122575 non-null float64
4   precip                122565 non-null float64
5   precipprob            122558 non-null float64
6   preciptype            11277 non-null  object

```

7	snow	89810	non-null	float64
8	snowdepth	89279	non-null	float64
9	windspeed	122583	non-null	float64
10	winddir	122567	non-null	float64
11	sealevelpressure	122319	non-null	float64
12	cloudcover	122556	non-null	float64
13	solarradiation	122514	non-null	float64
14	solarenergy	122480	non-null	float64
15	uvindex	122486	non-null	float64

dtypes: float64(14), object(2)

memory usage: 15.0+ MB

None

Print dimensions of the dataset

```
print("Dataset dimensions:", f_brighton_df.shape)
```

Dataset dimensions: (122844, 16)

Print first 2 rows of the dataset

```
print("First few rows of the dataset:")
```

```
print(f_brighton_df.head(2))
```

First few rows of the dataset:

	datetime	temp	dew	humidity	precip	precipprob
precip_type \						
0	2010-03-22T00:00:00	4.4	4.3	99.49	0.0	0.0
NaN						
1	2010-03-22T01:00:00	4.8	4.8	99.82	0.0	0.0
NaN						

	snow	snowdepth	windspeed	winddir	sealevelpressure
cloudcover \					
0	0.0	0.0	5.7	185.0	1021.1
					40.0
1	0.0	0.0	5.9	162.0	1020.9
					91.9

	solarradiation	solarenergy	uvindex
0	0.0	0.0	0.0
1	0.0	0.0	0.0

Summary

```
print("Summary:")
```

```
print(f_brighton_df.describe())
```

Summary:

	temp	dew	humidity	precip \
count	122590.000000	122568.000000	122575.000000	122565.000000
mean	11.059431	8.047911	82.880536	0.081954
std	5.654195	5.136014	12.511270	0.660720
min	-9.100000	-11.600000	24.340000	0.000000

25%	7.200000	4.500000	75.560000	0.000000
50%	11.000000	8.400000	85.700000	0.000000
75%	15.300000	12.000000	92.790000	0.000000
max	33.300000	20.200000	100.000000	32.385000

	precipprob	snow	snowdepth	windspeed \
count	122558.000000	89810.000000	89279.000000	122583.000000
mean	8.883141	0.000408	0.028941	15.938294
std	28.450141	0.034549	0.441740	8.903724
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	9.400000
50%	0.000000	0.000000	0.000000	14.400000
75%	0.000000	0.000000	0.000000	21.300000
max	100.000000	7.870000	96.000000	72.200000

	winddir	sealevelpressure	cloudcover	solarradiation
\				
count	122567.000000	122319.000000	122556.000000	122514.000000
mean	196.561486	1015.525463	60.841873	138.575974
std	106.273116	10.519485	31.879025	220.733866
min	0.700000	955.000000	0.000000	0.000000
25%	113.000000	1009.600000	36.000000	0.000000
50%	223.000000	1016.400000	68.400000	9.000000
75%	267.000000	1022.500000	89.800000	201.000000
max	360.000000	1049.300000	100.000000	1150.000000

	solarenergy	uvindex
count	122480.000000	122486.000000
mean	0.498450	1.368646
std	0.795407	2.225490
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.700000	2.000000
max	4.100000	10.000000

```
# Counting for missing values
print("Missing values:")
print(f_brighton_df.isnull().sum())
```

```
Missing values:
datetime      0
temp         254
```

```

dew                276
humidity           269
precip             279
precipprob         286
precipdtype        111567
snow              33034
snowdepth          33565
windspeed          261
winddir            277
sealevelpressure   525
cloudcover         288
solarradiation     330
solarenergy        364
uvindex            358
dtype: int64

# Info about categorical variables
print("Categories & frequencies for categorical variables:")
for col in f_brighton_df.select_dtypes(include='object').columns:
    print(f_brighton_df[col].value_counts())

Categories & frequencies for categorical variables:
2020-10-25T01:00:00    2
2014-10-26T01:00:00    2
2017-10-29T01:00:00    2
2022-10-30T01:00:00    2
2023-10-29T01:00:00    2
..
2014-09-03T02:00:00    1
2014-09-03T01:00:00    1
2014-09-03T00:00:00    1
2014-09-02T23:00:00    1
2023-10-19T23:00:00    1
Name: datetime, Length: 122830, dtype: int64
rain                10903
rain,snow           344
snow                30
Name: precipdtype, dtype: int64

```

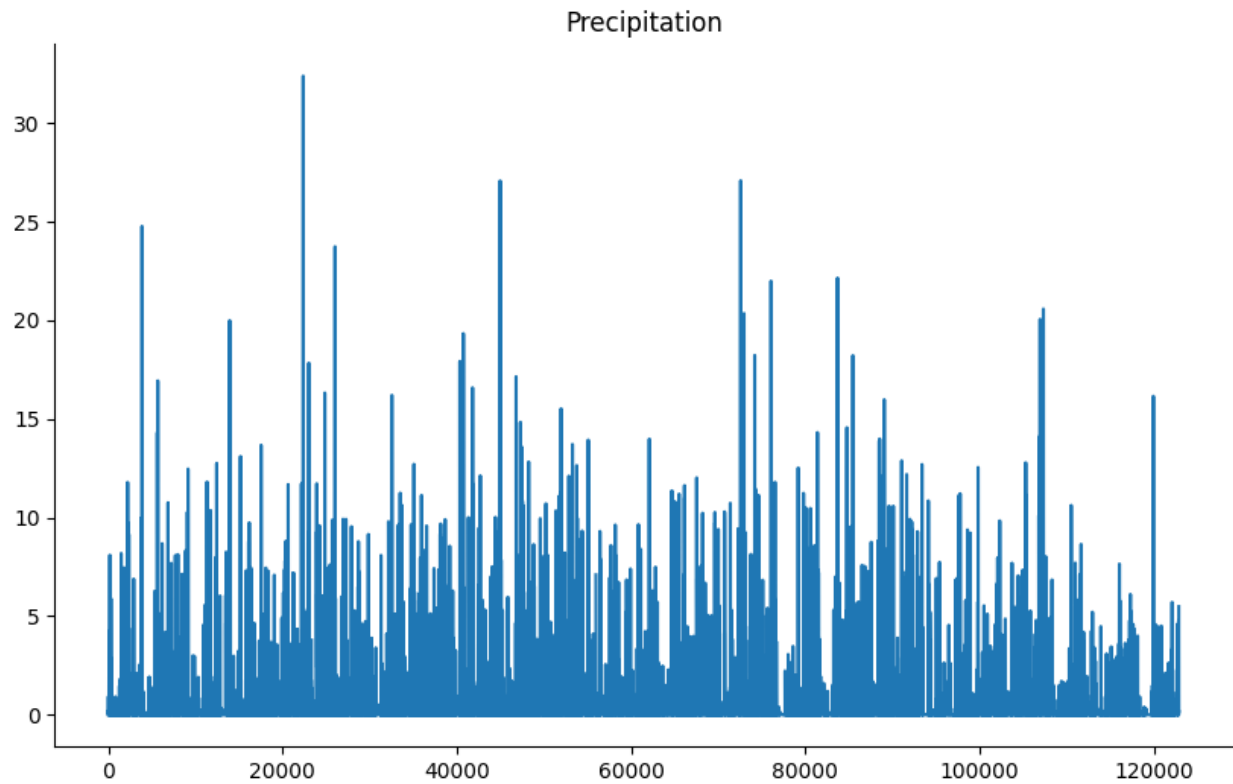
Graphs for EDA

```

# @title Precipitation


from matplotlib import pyplot as plt
f_brighton_df['precip'].plot(kind='line', figsize=(10, 6),
title='Precipitation')
plt.gca().spines[['top', 'right']].set_visible(False)

```




```
# @title Temperature
```

```
from matplotlib import pyplot as plt
f_brighton_df['temp'].plot(kind='line', figsize=(8, 4), title='temp')
plt.gca().spines[['top', 'right']].set_visible(False)
```

```
# @title Humidity

from matplotlib import pyplot as plt
f_brighton_df['humidity'].plot(kind='line', figsize=(8, 4),
title='humidity')
plt.gca().spines[['top', 'right']].set_visible(False)
```




```
# @title DateTime VS Precipitation

from matplotlib import pyplot as plt
import seaborn as sns
def _plot_series(series, series_name, series_index=0):
    from matplotlib import pyplot as plt
    import seaborn as sns
    palette = list(sns.palettes.mpl_palette('rocket'))
    xs = series['datetime']
    ys = series['precip']

    plt.plot(xs, ys, label=series_name, color=palette[series_index %
len(palette)])

fig, ax = plt.subplots(figsize=(10, 5.2), layout='constrained')
df_sorted = f_brighton_df.sort_values('datetime', ascending=True)
for i, (series_name, series) in
enumerate(df_sorted.groupby('preciptype')):
    _plot_series(series, series_name, i)
    fig.legend(title='preciptype', bbox_to_anchor=(1, 1), loc='upper
left')
sns.despine(fig=fig, ax=ax)
plt.xlabel('Datetime')
_ = plt.ylabel('Precipitation')
```



```
# @title Time VS Dew

from matplotlib import pyplot as plt
import seaborn as sns
def _plot_series(series, series_name, series_index=0):
    from matplotlib import pyplot as plt
```


```

import seaborn as sns
palette = list(sns.palettes.mpl_palette('flare'))
xs = series['datetime']
ys = series['dew']

plt.plot(xs, ys, label=series_name, color=palette[series_index %
len(palette)])

fig, ax = plt.subplots(figsize=(10, 5.2), layout='constrained')
df_sorted = f_brighton_df.sort_values('datetime', ascending=True)
for i, (series_name, series) in
enumerate(df_sorted.groupby('preciptype')):
    _plot_series(series, series_name, i)
    fig.legend(title='preciptype', bbox_to_anchor=(1, 1), loc='upper
left')
sns.despine(fig=fig, ax=ax)
plt.xlabel('Time')
_ = plt.ylabel('Dew')

```



```

# @title Time VS Humidity


from matplotlib import pyplot as plt
import seaborn as sns
def _plot_series(series, series_name, series_index=0):
    from matplotlib import pyplot as plt
    import seaborn as sns
    palette = list(sns.palettes.mpl_palette('crest'))
    xs = series['datetime']
    ys = series['humidity']

    plt.plot(xs, ys, label=series_name, color=palette[series_index %

```

```
len(palette)])

fig, ax = plt.subplots(figsize=(10, 5.2), layout='constrained')
df_sorted = f_brighton_df.sort_values('datetime', ascending=True)
for i, (series_name, series) in
    enumerate(df_sorted.groupby('preciptype')):
        _plot_series(series, series_name, i)
        fig.legend(title='preciptype', bbox_to_anchor=(1, 1), loc='upper
left')
sns.despine(fig=fig, ax=ax)
plt.xlabel('Time')
_ = plt.ylabel('Humidity')
```



```
# @title Time vs Temperature

from matplotlib import pyplot as plt
import seaborn as sns
def _plot_series(series, series_name, series_index=0):
    from matplotlib import pyplot as plt
    import seaborn as sns
    palette = list(sns.palettes.mpl_palette('Dark2'))
    xs = series['datetime']
    ys = series['temp']


    plt.plot(xs, ys, label=series_name, color=palette[series_index %
len(palette)])

fig, ax = plt.subplots(figsize=(10, 5.2), layout='constrained')
df_sorted = f_brighton_df.sort_values('datetime', ascending=True)
for i, (series_name, series) in
    enumerate(df_sorted.groupby('preciptype')):
```

```

    _plot_series(series, series_name, i)
    fig.legend(title='precipdtype', bbox_to_anchor=(1, 1), loc='upper
left')
sns.despine(fig=fig, ax=ax)
plt.xlabel('Time')
_ = plt.ylabel('Temperature')

```




```

# @title Precipitaion VS Precipitation Probability


from matplotlib import pyplot as plt
sns.boxplot(x=f_brighton_df['precip'], y=f_brighton_df['precipprob'])

<Axes: xlabel='precip', ylabel='precipprob'>

```

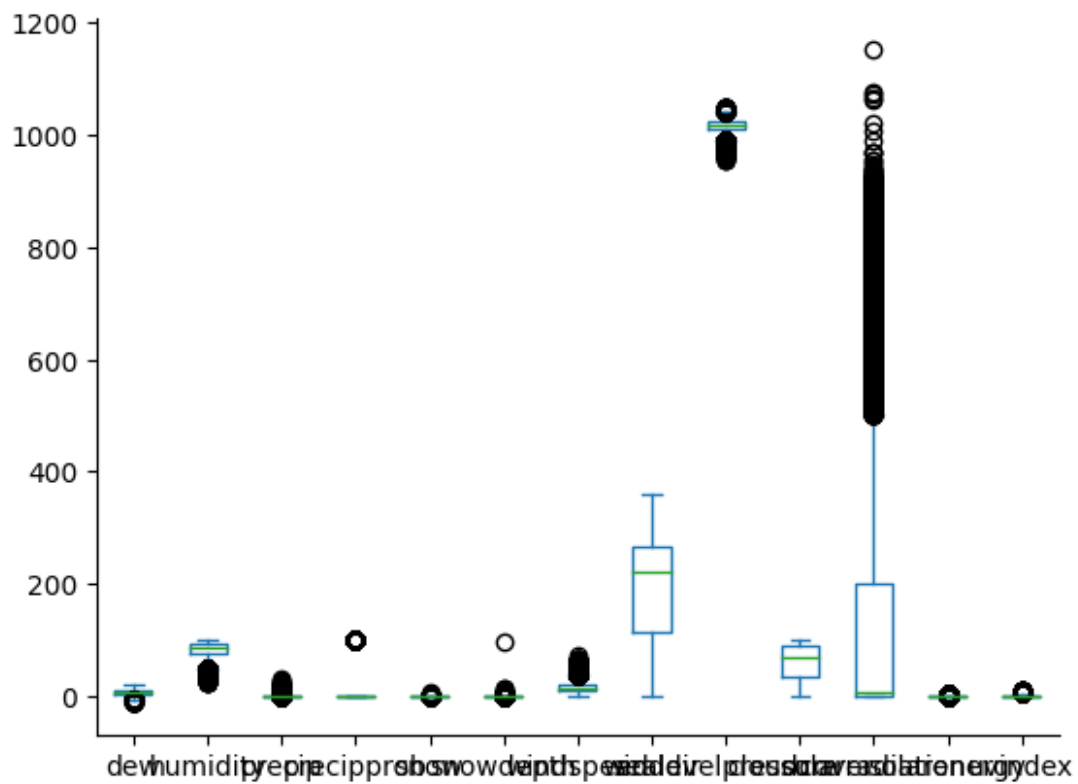


```
# @title Precipitaion VS Precipitation Probability
from matplotlib import pyplot as plt
f_brighton_df.plot(kind='scatter', x='precip', y='precipprob', s=32,
alpha=.8)
plt.gca().spines[['top', 'right']].set_visible(False)
```




```
# @title Temperature vs Dew
```

```
from matplotlib import pyplot as plt
f_brighton_df.plot(kind='box', x='temp')
plt.gca().spines[['top', 'right']].set_visible(False)
```




```
# @title Dew VS Humidity
```

```
from matplotlib import pyplot as plt
f_brighton_df.plot(kind='box', x='dew', y='humidity')
plt.gca().spines[['top', 'right']].set_visible(False)
```

```
# @title Humidity VS Precipitation


from matplotlib import pyplot as plt
f_brighton_df.plot(kind='scatter', x='humidity', y='precip', s=32,
alpha=.8)
plt.gca().spines[['top', 'right',]].set_visible(False)
```



###Exploring Important Columns


```
# @title Precipitation Type
```

```
from matplotlib import pyplot as plt
import seaborn as sns
f_brighton_df.groupby('preciptype').size().plot(kind='barh',
color=sns.palettes.mpl_palette('cubehelix'))
plt.gca().spines[['top', 'right',]].set_visible(False)
```



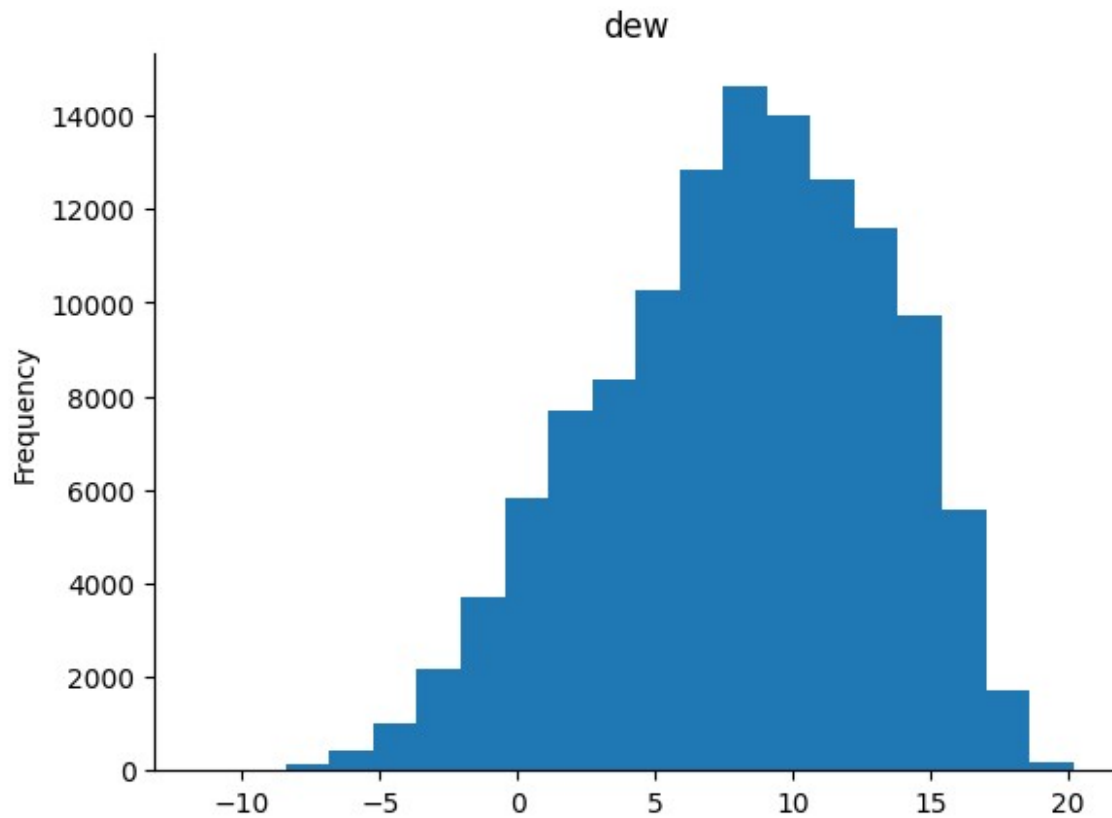
```
# @title Temperature
```

```
from matplotlib import pyplot as plt
f_brighton_df['temp'].plot(kind='hist', bins=20, title='temp')
plt.gca().spines[['top', 'right']].set_visible(False)
```




```
# @title Dew

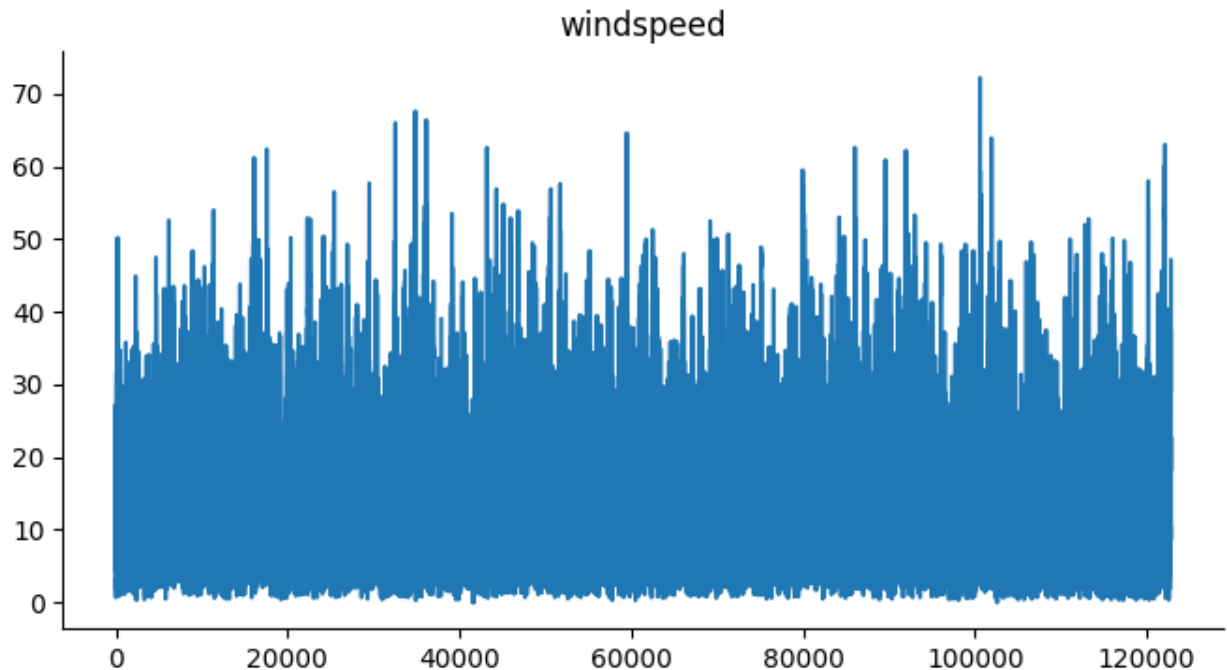
from matplotlib import pyplot as plt
f_brighton_df['dew'].plot(kind='hist', bins=20, title='dew')
plt.gca().spines[['top', 'right']].set_visible(False)
```



```
# @title Humidity  
from matplotlib import pyplot as plt  
f_brighton_df['humidity'].plot(kind='hist', bins=20, title='humidity')  
plt.gca().spines[['top', 'right',]].set_visible(False)
```



```
# @title Windspeed
from matplotlib import pyplot as plt
f_brighton_df['windspeed'].plot(kind='line', figsize=(8, 4),
title='windspeed')
plt.gca().spines[['top', 'right']].set_visible(False)
```



##Time Series Analysis

```
months_of_the_year = ['January', 'February', 'March', 'April', 'May',
                       'June', 'July', 'August', 'September', 'October', 'November',
                       'December']
```

```
f_brighton_df.columns
```

```
Index(['datetime', 'temp', 'dew', 'humidity', 'precip', 'precipprob',
       'preciptype', 'snow', 'snowdepth', 'windspeed', 'winddir',
       'sealevelpressure', 'cloudcover', 'solarradiation',
       'solarenergy',
       'uvindex'],
      dtype='object')
```

```
# Date will be our index. Let's convert it to a datetime type
```

```
f_brighton_df['datetime'] = pd.to_datetime(f_brighton_df['datetime'],
                                           dayfirst=True)
```

```
f_brighton_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 122844 entries, 0 to 122843
Data columns (total 16 columns):
```

#	Column	Non-Null Count	Dtype
0	datetime	122844 non-null	datetime64[ns]
1	temp	122590 non-null	float64
2	dew	122568 non-null	float64
3	humidity	122575 non-null	float64

```

4    precip          122565 non-null float64
5    precipprob      122558 non-null float64
6    preciptype       11277 non-null  object
7    snow             89810 non-null float64
8    snowdepth        89279 non-null float64
9    windspeed        122583 non-null float64
10   winddir          122567 non-null float64
11   sealevelpressure 122319 non-null float64
12   cloudcover        122556 non-null float64
13   solarradiation    122514 non-null float64
14   solarenergy       122480 non-null float64
15   uvindex          122486 non-null float64
dtypes: datetime64[ns](1), float64(14), object(1)
memory usage: 15.0+ MB

```

```

print(f_brighton_df['datetime'].min(),
f_brighton_df['datetime'].max())

```

```

2010-01-01 00:00:00 2024-01-06 11:00:00

```

Let's say we want to create extra columns: month, year, and week of the year

```

df = f_brighton_df.copy()
df['month'] = df['datetime'].dt.month
df['year'] = df['datetime'].dt.year
df['week_of_year'] = df['datetime'].dt.isocalendar().week
df

```

```

{"type": "dataframe", "variable_name": "df"}

```

Let's make the date column the index of the dataframe for easier slicing

```

df.set_index('datetime', inplace=True) # note we can only run this
once, as it will delete the 'date' column.
df.head()

```

```

{"type": "dataframe", "variable_name": "df"}

```

```

print(df.isna().any())
df=df.dropna()
print(df.isna().any())

```

```

temp          True
dew            True
humidity       True
precip         True
precipprob     True
preciptype     True
snow           True
snowdepth      True
windspeed      True


```



```
winddir      True
sealevelpressure  True
cloudcover   True
solarradiation  True
solarenergy   True
uvindex       True
month         False
year          False
week_of_year  False
dtype: bool
temp         False
dew          False
humidity     False
precip       False
precipprob   False
preciptype   False
snow         False
snowdepth    False
windspeed    False
winddir      False
sealevelpressure  False
cloudcover   False
solarradiation  False
solarenergy   False
uvindex       False
month         False
year          False
week_of_year  False
dtype: bool
```


Let's plot the data. For now we're only going to work on temp_avg, so let's have a look

```
plt.figure(figsize=(16,8))
plt.plot(df.index, df['temp'])
plt.xlabel('Year')
plt.ylabel('temperature (C)')
Text(0, 0.5, 'temperature (C)')
```



```
# Let's zoom in to 2014-2017
df_chunk = df.loc['2014-12':'2017-01'] # since the date is an index,
we can use it to filter our data

plt.figure(figsize=(10, 8))
plt.plot(df_chunk.index, df_chunk['temp'])
plt.xticks(rotation=90)
plt.xlabel('Date')
plt.ylabel('Average temperature (C)')
```



```
print(df_chunk.reindex(pd.date_range('2014-12', '2017-01')).isnull().all(1).sum()) # 33 days missing
df_chunk.reindex(pd.date_range('2014-12', '2017-01')).isnull().all(1)
```

752

2014-12-01	True
2014-12-02	True
2014-12-03	True
2014-12-04	True
2014-12-05	True
...	
2016-12-28	True
2016-12-29	True
2016-12-30	True
2016-12-31	False

2017-01-01 True
Freq: D, Length: 763, dtype: bool

```
df[df.index.duplicated(keep=False)].head(20)
```

```
{
  "summary": {
    "\n  \"name\": \"df[df\", \n  \"rows\": 2, \n  \"fields\": [
      \n    {
      \n      \"column\": \"temp\", \n      \"properties\": {
        \"dtype\": \"number\", \n        \"std\": 0.14142135623731025, \n        \"min\": 11.2, \n        \"max\": 11.4, \n        \"samples\": [
          11.2, \n          11.4 \n        ], \n        \"num_unique_values\": 2, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      } \n    }, \n    {
      \n      \"column\": \"dew\", \n      \"properties\": {
        \"dtype\": \"number\", \n        \"std\": 0.2828427124746193, \n        \"min\": 10.1, \n        \"max\": 10.5, \n        \"samples\": [
          10.1, \n          10.5 \n        ], \n        \"num_unique_values\": 2, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      } \n    }, \n    {
      \n      \"column\": \"humidity\", \n      \"properties\": {
        \"dtype\": \"number\", \n        \"std\": 0.5374011537017798, \n        \"min\": 93.11, \n        \"max\": 93.87, \n        \"samples\": [
          93.11, \n          93.87 \n        ], \n        \"num_unique_values\": 2, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      } \n    }, \n    {
      \n      \"column\": \"precip\", \n      \"properties\": {
        \"dtype\": \"number\", \n        \"std\": 1.233194226389339, \n        \"min\": 0.678, \n        \"max\": 2.422, \n        \"samples\": [
          2.422, \n          0.678 \n        ], \n        \"num_unique_values\": 2, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      } \n    }, \n    {
      \n      \"column\": \"precipprob\", \n      \"properties\": {
        \"dtype\": \"number\", \n        \"std\": 0.0, \n        \"min\": 100.0, \n        \"max\": 100.0, \n        \"samples\": [
          100.0 \n        ], \n        \"num_unique_values\": 1, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      } \n    }, \n    {
      \n      \"column\": \"preciptype\", \n      \"properties\": {
        \"dtype\": \"string\", \n        \"samples\": [
          \"rain\" \n        ], \n        \"num_unique_values\": 1, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      } \n    }, \n    {
      \n      \"column\": \"snow\", \n      \"properties\": {
        \"dtype\": \"number\", \n        \"std\": 0.0, \n        \"min\": 0.0, \n        \"max\": 0.0, \n        \"samples\": [
          0.0 \n        ], \n        \"num_unique_values\": 1, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      } \n    }, \n    {
      \n      \"column\": \"snowdepth\", \n      \"properties\": {
        \"dtype\": \"number\", \n        \"std\": 0.0, \n        \"min\": 0.0, \n        \"max\": 0.0, \n        \"samples\": [
          0.0 \n        ], \n        \"num_unique_values\": 1, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      } \n    }, \n    {
      \n      \"column\": \"windspeed\", \n      \"properties\": {
        \"dtype\": \"number\", \n        \"std\": 7.424621202458749, \n        \"min\": 11.3, \n        \"max\": 21.8, \n        \"samples\": [
          21.8 \n        ]
      }
    }
  ]
}
```

```
[,\n      \nnum_unique_values\": 2,\n\"semantic_type\": \"\",,\n      },\n      {\n        \ncolumn\": \"winddir\",,\n        \nproperties\": {\n          \ndtype\": \"number\",,\n          \nstd\": 3.5355339059327378,\n          \nmin\": 245.0,\n          \nmax\": 250.0,\n          \nsamples\": [\n            245.0\n          ],\n        }\n      },\n      {\n        \ncolumn\": \nsealevelpressure\",,\n        \nproperties\": {\n          \ndtype\": \"number\",,\n          \nstd\": 0.6363961030678768,\n          \nmin\": 983.7,\n          \nmax\": 984.6,\n          \nsamples\": [\n            984.6\n          ],\n        }\n      },\n      {\n        \ncolumn\": \ncloudcover\",,\n        \nproperties\": {\n          \ndtype\": \"number\",,\n          \nstd\": 6.505382386916239,\n          \nmin\": 86.6,\n          \nmax\": 95.8,\n          \nsamples\": [\n            86.6\n          ],\n        }\n      },\n      {\n        \ncolumn\": \nsolarradiation\",,\n        \nproperties\": {\n          \ndtype\": \"number\",,\n          \nstd\": 0.0,\n          \nmin\": 0.0,\n          \nmax\": 0.0,\n          \nsamples\": [\n            0.0\n          ],\n        }\n      },\n      {\n        \ncolumn\": \nsolarenergy\",,\n        \nproperties\": {\n          \ndtype\": \"number\",,\n          \nstd\": 0.0,\n          \nmin\": 0.0,\n          \nmax\": 0.0,\n          \nsamples\": [\n            0.0\n          ],\n        }\n      },\n      {\n        \ncolumn\": \nuvindex\",,\n        \nproperties\": {\n          \ndtype\": \"number\",,\n          \nstd\": 0.0,\n          \nmin\": 0.0,\n          \nmax\": 0.0,\n          \nsamples\": [\n            0.0\n          ],\n        }\n      },\n      {\n        \ncolumn\": \nmonth\",,\n        \nproperties\": {\n          \ndtype\": \"number\",,\n          \nstd\": 0,\n          \nmin\": 10,\n          \nmax\": 10,\n          \nsamples\": [\n            10\n          ],\n        }\n      },\n      {\n        \ncolumn\": \nyear\",,\n        \nproperties\": {\n          \ndtype\": \"number\",,\n          \nstd\": 0,\n          \nmin\": 2023,\n          \nmax\": 2023,\n          \nsamples\": [\n            2023\n          ],\n        }\n      },\n      {\n        \ncolumn\": \nweek_of_year\",,\n        \nproperties\": {\n          \ndtype\": \"UInt32\",,\n          \nsamples\": [\n            43\n          ],\n        }\n      },\n      {\n        \ncolumn\": \n\n      },\n    ]\n  },\n  \"type\": \"dataframe\"}
```

```
# Let's keep the first one only - in practice this would require more careful analysis!
```

```
df = df[~df.index.duplicated(keep='first')]
len(df)
```

```
8535
```

```
# Now we can reindex -- this is where the original error about duplicates was
```

```
df = df.reindex(pd.date_range(df.index[0], df.index[-1]))
print(len(df))
```

```
4960
```


```
# Now we should have missing values
```

```
print(df.isna().sum())
```

temp	4191
dew	4191
humidity	4191
precip	4191
precipprob	4191
preciptype	4191
snow	4191
snowdepth	4191
windspeed	4191
winddir	4191
sealevelpressure	4191
cloudcover	4191
solarradiation	4191
solarenergy	4191
uvindex	4191
month	4191
year	4191
week_of_year	4191
dtype:	int64

```
df_chunk = df.loc['2017-12-15':'2018-01-15'] # since the date is an index, we can use it to filter our data
```

```
plt.figure(figsize=(10, 8))
plt.plot(df_chunk.index, df_chunk['temp'])
plt.xticks(rotation=90)
plt.xlabel('Date')
plt.ylabel('temperature (C)')
# The missing values are clearly visible now!
```



```
df2 = df_chunk.copy()
df2 = df2.loc[:, 'temp'].to_frame()
df2


{"summary":{"\n  \"name\": \"df2\",\n  \"rows\": 32,\n  \"fields\": [\n    {\n      \"column\": \"temp\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 2.5582622778661577,\n        \"min\": 2.4,\n        \"max\": 11.7,\n        \"samples\": [\n          7.1,\n          3.7,\n          3.4\n        ],\n        \"num_unique_values\": 15,\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ]\n}, \"type\": \"dataframe\", \"variable_name\": \"df2\"}

#Forward Fill
df2['ffill'] = df2['temp'].ffill()
# Backward Fill
df2['bfill'] = df2['temp'].bfill()
```

```
# Mean Value Fill
df2['meanfill'] = df2['temp'].fillna(df['temp'].mean()) # Note that
we're using the mean of df, not of df2
# Fill with 0s
df2['zerofill'] = df2['temp'].fillna(0)

# Plot
fig, ax = plt.subplots(figsize=(10,8))

plt.plot(df2.index, df2['ffill'], label='ffill', linestyle='--',
color='red')
plt.plot(df2.index, df2['bfill'], label='bfill', linestyle='-.',
color='green')
plt.plot(df2.index, df2['meanfill'], label='mean', linestyle=':',
color='black')
plt.plot(df2.index, df2['zerofill'], linestyle='--', color='purple',
label='zero')
plt.plot(df2.index, df2['temp'], color='blue', label='Original')
plt.legend()
plt.ylabel('Temperature')
plt.ylim(-1, 26)
_=plt.title('Forward, backward, zero, and mean value fill')
```

Try different ways to fill the data - more advanced: interpolation

```
df2['linear_interp'] = df2['temp'].interpolate(method='linear')
df2['nearest_interp'] = df2['temp'].interpolate(method='nearest')
df2['spline_interp'] = df2['temp'].interpolate(method='spline',
order=2)
df2['polynomial_interp'] =
df2['temp'].interpolate(method="polynomial", order=3)
```

Plot

```
fig, ax = plt.subplots(figsize=(10,8))
```


```
plt.plot(df2.index, df2['linear_interp'], linestyle='--', color='red',
label='linear')
plt.plot(df2.index, df2['nearest_interp'], linestyle='-.',
color='green', label='nearest')
plt.plot(df2.index, df2['spline_interp'], linestyle=':',
color='black', label='spline')
plt.plot(df2.index, df2['polynomial_interp'], linestyle='--',
```

```

color='purple', label='polynomial')
plt.plot(df2.index, df2['temp'], label='Original', color='blue')

plt.legend()
plt.ylabel('Temperature')
plt.ylim(-1, 26)
_=plt.title('Interpolation')


```



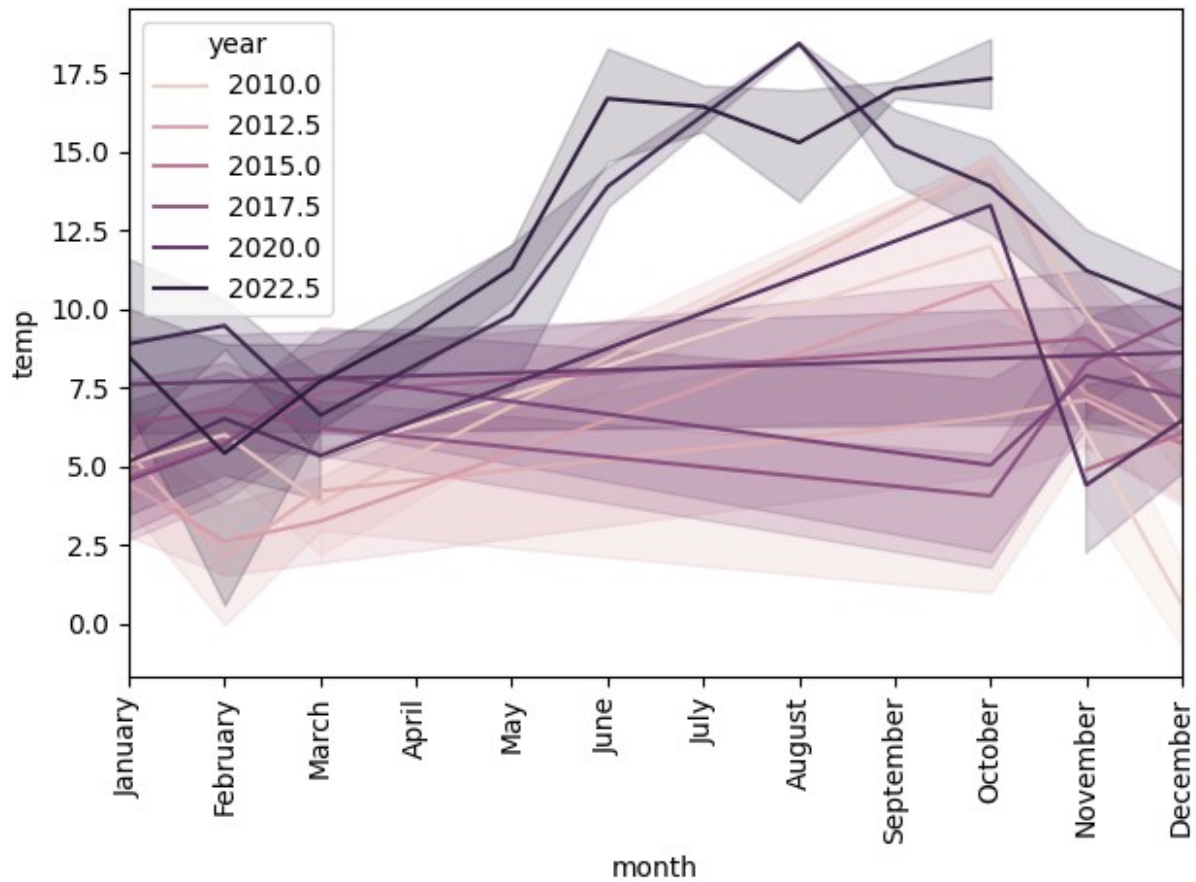
```

plt.figure(figsize=(16, 8))
_=sns.boxplot(x='year', y='temp', data=df)
_=plt.tight_layout()


```



```
# Visualise trends across years
sns.lineplot(x='month', y='temp', data=df, hue='year')
_ = plt.xticks(np.arange(1, 13), months_of_the_year, rotation=90)
_ = plt.xlim(1, 12) # limit x-axis
_ = plt.tight_layout()
```




```
# Visualise trends across years
fix, ax = plt.subplots(2, 1, sharex=True, figsize=(20,15))
sns.boxplot(x='month', y='temp', data=df, ax=ax[0]) # top plot
sns.boxplot(x='month', y='temp', data=df, hue='year', ax=ax[1]) #
bottom plot
ax[1].set_xticks(np.arange(0, 12), months_of_the_year, fontsize=14)
plt.tight_layout()
```




```
data_ds = df['temp'].resample('M').mean().ffill().to_frame() # one
value per month
data_ds
```

```
{"summary":{"\n  \"name\": \"data_ds\", \n  \"rows\": 164, \n  \"fields\": [\n    {\n      \"column\": \"temp\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 3.0209873807308147, \n        \"min\": 0.6166666666666666, \n        \"max\": 18.45, \n        \"samples\": [\n          8.216666666666667, \n          5.366666666666667, \n          6.463636363636363\n        ], \n        \"num_unique_values\": 70, \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    }\n  ], \n  \"type\": \"dataframe\", \n  \"variable_name\": \"data_ds\"}
```

```
# Try decomposition on the resampled dataset
from statsmodels.tsa.seasonal import seasonal_decompose, STL
decomposition = STL(data_ds['temp']).fit()
fig = decomposition.plot()
fig.set_size_inches(12,8)
fig.axes[1].grid()
```



```
# Try decomposition on the resampled dataset, using only the full
years
decomposition = STL(data_ds.loc['2017':'2024', 'temp']).fit()
fig = decomposition.plot()
fig.set_size_inches(12,8)
fig.axes[1].grid()
```



```
# Statistical test for stationarity: Augmented Dickey-Fuller (ADF) test
```

```
adf_result = adfuller(data_ds['temp'])
print('ADF Statistic %.2f:' % adf_result[0])
print('ADF p-value: %.4f:' % adf_result[1])
# p-value < 0.05 ==> timeseries does not have a unit root
```

```
ADF Statistic -0.13:
ADF p-value: 0.9465:
```


```
adf_result = adfuller(data_ds.loc['2017':'2024', 'temp']) # ADF test on the full years only. Is there a trend?
```

```
print('ADF Statistic %.2f:' % adf_result[0])
print('ADF p-value: %.4f:' % adf_result[1])
```


```
ADF Statistic 0.10:
ADF p-value: 0.9659:
```

```
# Autocorrelation (can help us with modelling later)
```

```
fig, ax = plt.subplots(figsize=(16,8))
_ = plot_acf(data_ds['temp'], lags=48, ax=ax) # each lag is one month, so we're looking at 4 years worth of past data
_ = plt.xlabel('Lags (months)')
_ = plt.ylabel('Autocorrelation')
```



```
# Partial autocorrelation (can help us with modelling later)
fig, ax = plt.subplots(figsize=(16,8))
_ = plot_pacf(data_ds['temp'], ax=ax)
_ = plt.xlabel('Lags (months)')
_ = plt.ylabel('Partial autocorrelation')
```



```
# https://seaborn.pydata.org/examples/many\_pairwise\_correlations.html
# Compute the correlation matrix
```



```

corr = df.iloc[:, :-3].corr()

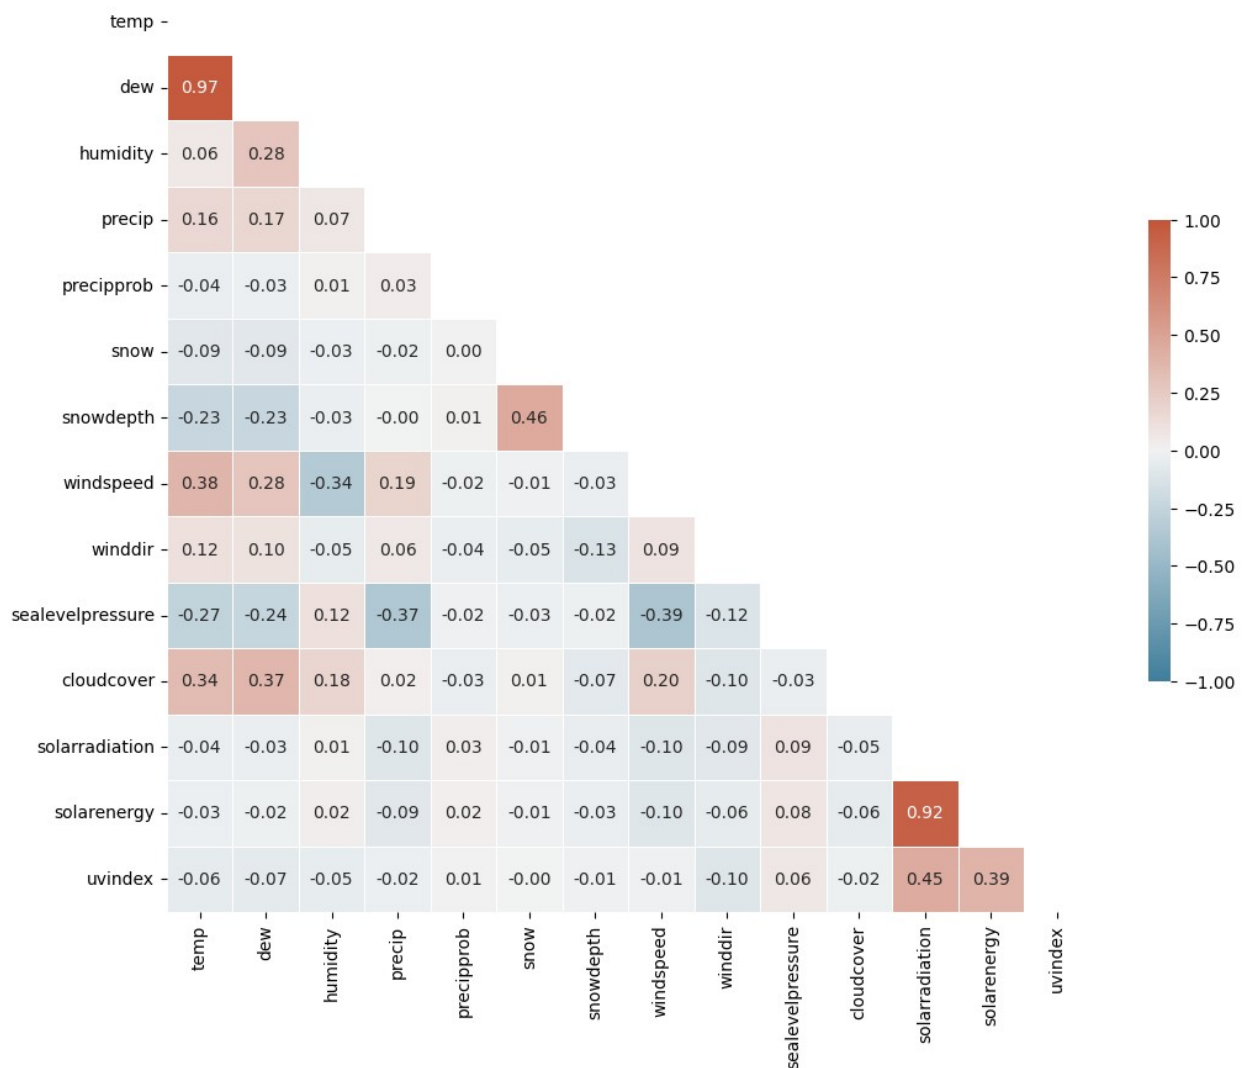
# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr, dtype=bool))

# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(11, 9))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(230, 20, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmin=-1, vmax=1, center=0,
            annot=True, square=True, linewidths=.5, cbar_kws={"shrink": .5},
            fmt='.2f')
plt.tight_layout()

```



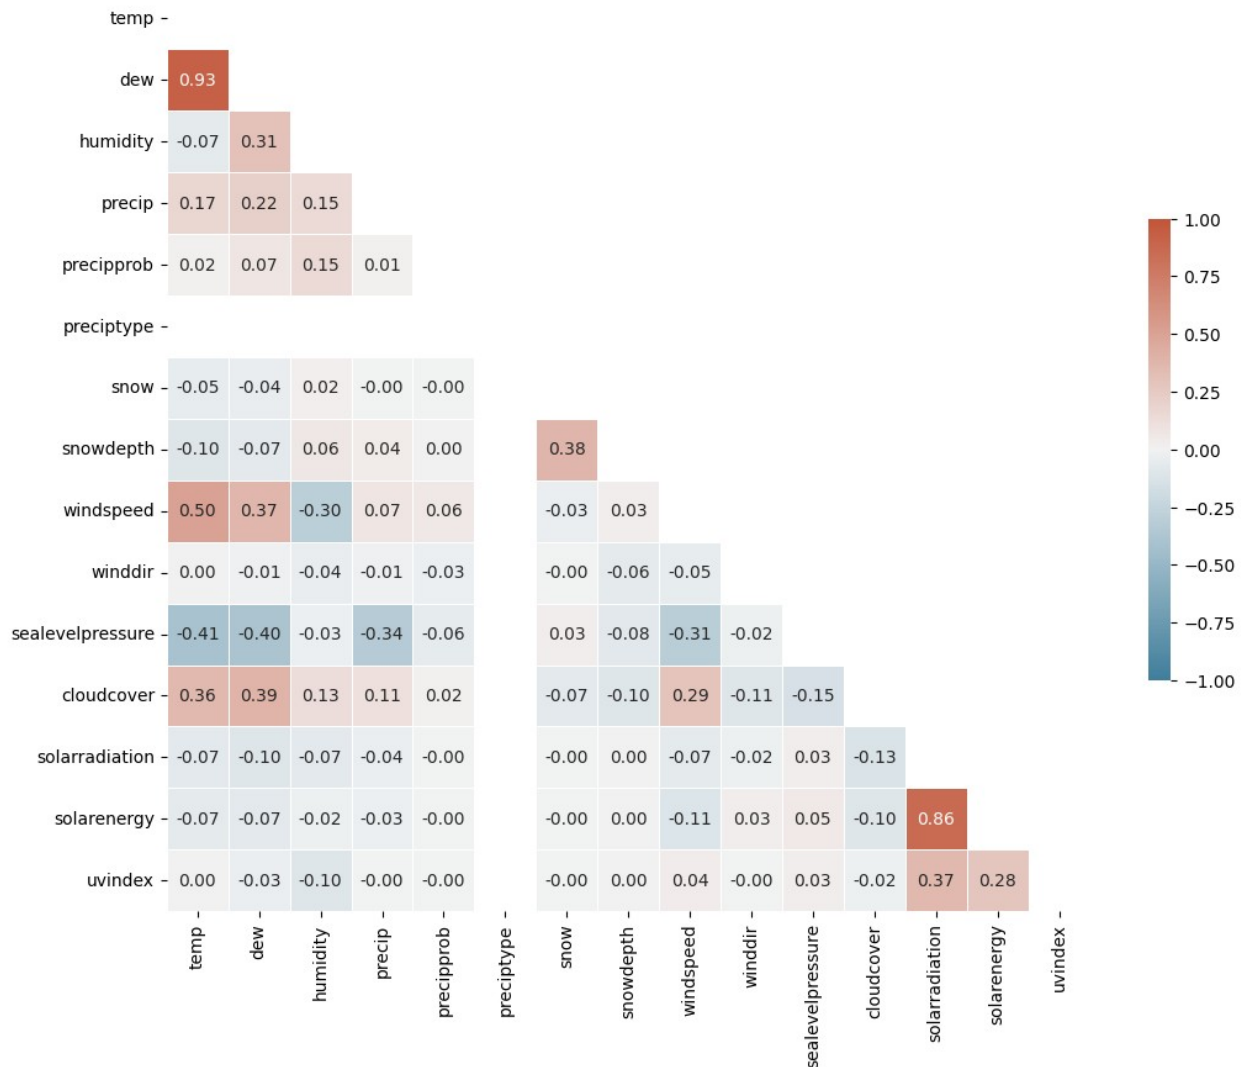
```
# Convert columns to numeric type if necessary
df_numeric = df.iloc[:, :-3].apply(pd.to_numeric, errors='coerce')

# Compute the correlation matrix
corr_diff = df_numeric.diff().corr()
# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr_diff, dtype=bool))

# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(11, 9))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(230, 20, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr_diff, mask=mask, cmap=cmap, vmin=-1, vmax=1,
            center=0, annot=True,
            square=True, linewidths=.5, cbar_kws={"shrink": .5},
            fmt='.2f')
plt.tight_layout()
plt.show()
```



```
df2 = df.iloc[:, :-3].copy()
df2['temp_avg_lag3'] = df2['temp'].shift(-3)
df2.head()
```

```
{
  "summary": {
    "name": "df2",
    "rows": 4960,
    "fields": [
      {
        "column": "temp",
        "properties": {
          "dtype": "number",
          "std": 4.502114672691266,
          "min": -5.8,
          "max": 18.6,
          "samples": [
            2.4,
            9.4,
            15.0
          ],
          "num_unique_values": 193,
          "semantic_type": "\"\"",
          "description": "\"\""}
        },
      {
        "column": "dew",
        "properties": {
          "dtype": "number",
          "std": 4.63943049276453,
          "min": -6.9,
          "max": 17.5,
          "samples": [
            8.5,
            -2.0,
            -1.6
          ],
          "num_unique_values": 198,
          "semantic_type": "\"\"",
          "description": "\"\""}
        },
      {
        "column": "humidity",
        "properties": {
          "dtype": "number",
          "std": 4.502114672691266,
          "min": -5.8,
          "max": 18.6,
          "samples": [
            2.4,
            9.4,
            15.0
          ],
          "num_unique_values": 193,
          "semantic_type": "\"\"",
          "description": "\"\""}
        }
      ]
    }
  }
}
```

```

{"dtype": "number", "std": 6.477026219315302, "min": 64.25, "max": 100.0, "samples": [92.36, 92.67, 92.7], "num_unique_values": 622, "semantic_type": "", "description": ""}, {"column": "precip", "properties": {"dtype": "number", "std": 2.2224425362941083, "min": 0.0, "max": 23.733, "samples": [1.228, 0.276, 1.609]}, "num_unique_values": 385, "semantic_type": "", "description": ""}, {"column": "precipprob", "properties": {"dtype": "number", "std": 8.042441332836129, "min": 0.0, "max": 100.0, "samples": [0.0, 100.0]}, "num_unique_values": 2, "semantic_type": "", "description": ""}, {"column": "preciptype", "properties": {"dtype": "category", "samples": ["rain", "rain,snow"], "num_unique_values": 3, "semantic_type": "", "description": ""}, {"column": "snow", "properties": {"dtype": "number", "std": 0.004903020264658105, "min": 0.0, "max": 0.13, "samples": [0.0, 0.04]}, "num_unique_values": 3, "semantic_type": "", "description": ""}, {"column": "snowdepth", "properties": {"dtype": "number", "std": 0.39570796414247517, "min": 0.0, "max": 5.0, "samples": [4.75, 2.0]}, "num_unique_values": 12, "semantic_type": "", "description": ""}, {"column": "windspeed", "properties": {"dtype": "number", "std": 10.368344360719101, "min": 0.8, "max": 66.0, "samples": [15.9, 30.8]}, "num_unique_values": 310, "semantic_type": "", "description": ""}, {"column": "winddir", "properties": {"dtype": "number", "std": 101.91043086560504, "min": 2.0, "max": 360.0, "samples": [197.0, 91.0]}, "num_unique_values": 293, "semantic_type": "", "description": ""}, {"column": "sealevelpressure", "properties": {"dtype": "number", "std": 13.091025102392456, "min": 968.1, "max": 1043.2, "samples": [1026.6, 987.2]}, "num_unique_values": 394, "semantic_type": "", "description": ""}, {"column": "cloudcover", "

```

```

{"properties": {"dtype": "number", "std": 27.910066073116568, "min": 0.0, "max": 100.0, "samples": [76.4, 50.3], "num_unique_values": 378, "semantic_type": "", "description": ""}, {"column": "solarradiation", "properties": {"dtype": "number", "std": 7.480077018670658, "min": 0.0, "max": 59.3, "samples": [9.0, 17.6], "num_unique_values": 97, "semantic_type": "", "description": ""}, {"column": "solarenergy", "properties": {"dtype": "number", "std": 0.030652888741679738, "min": 0.0, "max": 0.2, "samples": [0.1, 0.0], "num_unique_values": 3, "semantic_type": "", "description": ""}, {"column": "uvindex", "properties": {"dtype": "number", "std": 0.06237796929458344, "min": 0.0, "max": 1.0, "samples": [1.0, 0.0], "num_unique_values": 2, "semantic_type": "", "description": ""}, {"column": "temp_avg_lag3", "properties": {"dtype": "number", "std": 4.506497352455208, "min": -5.8, "max": 18.6, "samples": [6.8, 9.4], "num_unique_values": 193, "semantic_type": ""}}], "type": "dataframe", "variable_name": "df2"}

```

Let's see what happens if we do the differential operation again.

```

df2 = df.iloc[:, :-3].copy()
df2['temp_avg_lag3'] = df2['temp'].shift(-3)
corr2 = df2.corr()
# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr2, dtype=bool))

```

Set up the matplotlib figure

```
f, ax = plt.subplots(figsize=(11, 9))
```

Generate a custom diverging colormap

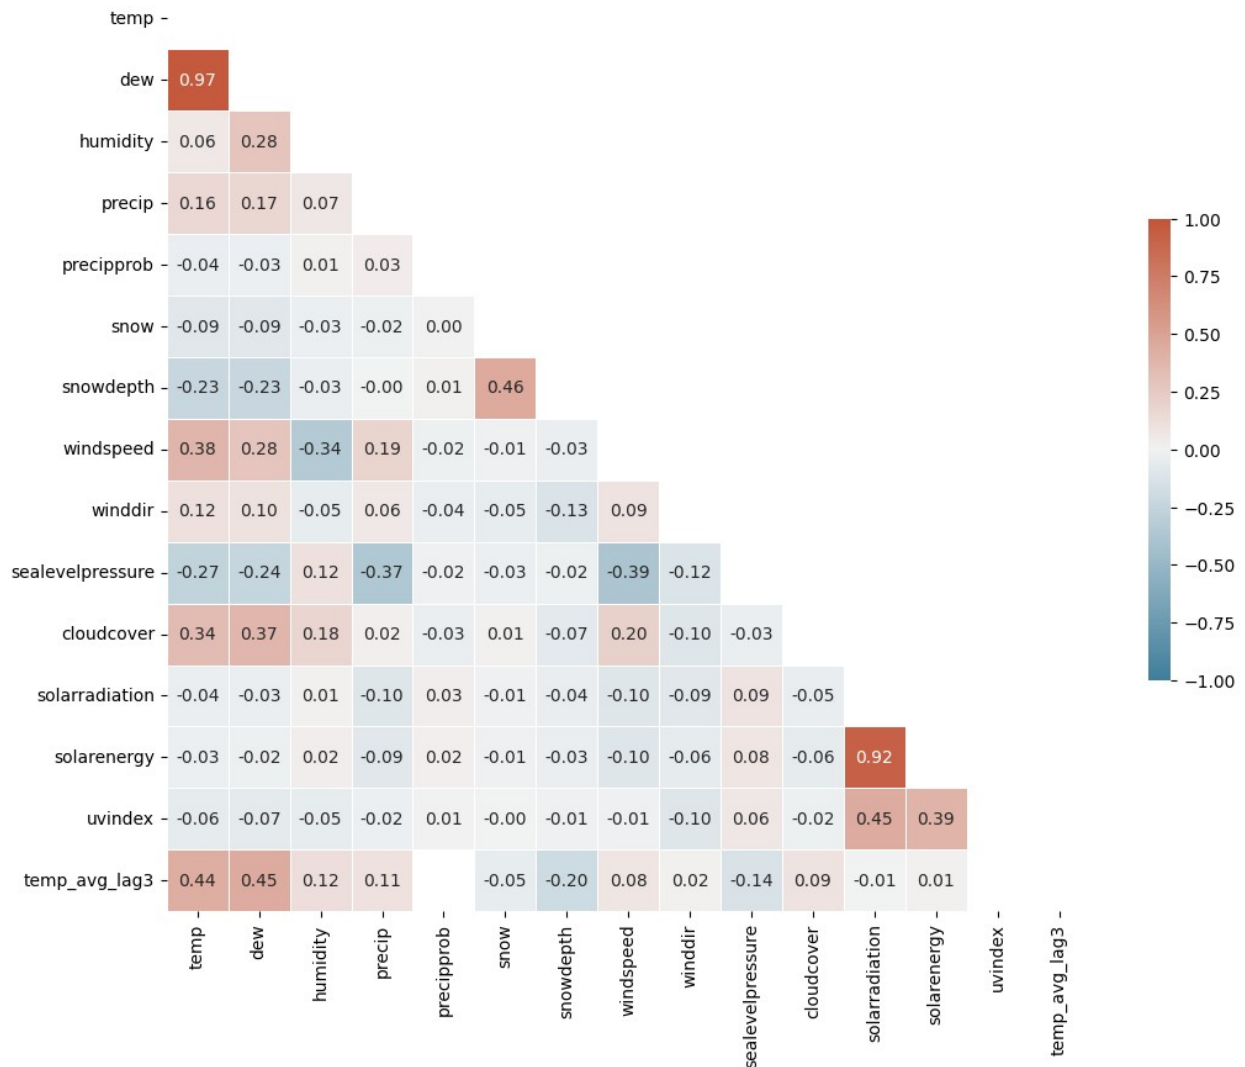
```
cmap = sns.diverging_palette(230, 20, as_cmap=True)
```

Draw the heatmap with the mask and correct aspect ratio

```

sns.heatmap(corr2, mask=mask, cmap=cmap, vmin=-1, vmax=1, center=0,
            annot=True, square=True, linewidths=.5, cbar_kws={"shrink": .5},
            fmt='.2f')
plt.tight_layout()

```



#Colchester Data Visualization

##Exploratory Data Analysis

```
# Specify the folder directory where the dataset is located.
df_path = GOOGLE_DRIVE_PATH

#list of all files in the path
file_list = [file for file in os.listdir(GOOGLE_DRIVE_PATH) if
file.endswith('.csv')]

brighton_df = []

# Iterate over each CSV file
for file_in file_list:
    # Construct the full path to the CSV file
```

```

file_path = os.path.join(df_path, file_)

# Check for file has an index column named '0'
index_column = pd.read_csv(file_path, nrows=1).columns[0] == '0'

# Change header value based on the value of the column
header = 1 if index_column else "infer"

current_dataframe = pd.read_csv(file_path, header=header)

# If the file name contains "Colchester", add its DataFrame to the
list
if file_.startswith('colchester'):
    brighton_df.append(current_dataframe)

# Concatenate all the datasets into one
f_colchester_df = pd.concat(brighton_df, ignore_index=True)

#information about the DataFrame
print(f_colchester_df.info())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 188024 entries, 0 to 188023
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   datetime              188024 non-null object
1   temp                  187606 non-null float64
2   dew                   187600 non-null float64
3   humidity              187602 non-null float64
4   precip                187437 non-null float64
5   precipprob            187590 non-null float64
6   preciptype            19448 non-null  object
7   snow                  186478 non-null float64
8   snowdepth             186429 non-null float64
9   windspeed             187627 non-null float64
10  winddir               187595 non-null float64
11  sealevelpressure      180462 non-null float64
12  cloudcover            184405 non-null float64
13  solarradiation        94894 non-null  float64
14  uvindex               94897 non-null  float64
15  solarenergy           94908 non-null  float64
dtypes: float64(14), object(2)
memory usage: 23.0+ MB
None

# Print dimensions of the dataset
print("Dataset dimensions:", f_colchester_df.shape)

Dataset dimensions: (188024, 16)

```

```
# Print first 2 rows of the dataset
print("First few rows of the dataset:")
print(f_colchester_df.head(2))
```

First few rows of the dataset:

	datetime	temp	dew	humidity	precip	precipprob
precip_type \						
0	2000-01-01T00:00:00	5.9	5.7	98.63	0.0	0.0
NaN						
1	2000-01-01T01:00:00	6.4	5.8	96.12	0.0	0.0
NaN						

	snow	snowdepth	windspeed	winddir	sealevelpressure	cloudcover
0	NaN	NaN	9.4	210.0	1020.6	NaN
1	NaN	NaN	15.1	233.0	1020.4	100.0

	solarradiation	uvindex	solarenergy
0	NaN	NaN	NaN
1	NaN	NaN	NaN

```
# Summary
print("Summary:")
print(f_colchester_df.describe())
```

Summary:

	temp	dew	humidity	precip
count	187606.000000	187600.000000	187602.000000	187437.000000
mean	10.503048	7.190720	81.596979	0.067773
std	6.004301	4.956358	14.085087	0.623144
min	-9.600000	-10.800000	22.430000	0.000000
25%	6.100000	3.600000	73.490000	0.000000
50%	10.300000	7.400000	85.480000	0.000000
75%	14.700000	10.900000	92.680000	0.000000
max	35.000000	21.500000	100.000000	84.324000

	precipprob	snow	snowdepth	windspeed
count	187590.000000	186478.000000	186429.000000	187627.000000
mean	10.276578	0.001008	0.062193	15.998126
std	30.357300	0.043173	0.602604	7.916611
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	10.100000
50%	0.000000	0.000000	0.000000	14.600000
75%	0.000000	0.000000	0.000000	20.700000
max	100.000000	9.100000	15.230000	74.700000

	winddir	sealevelpressure	cloudcover	solarradiation
count	187595.000000	180462.000000	184405.000000	94894.000000

mean	198.262143	1013.922618	59.510905	122.256657
std	95.653613	29.713385	31.985577	198.631750
min	0.000000	0.000000	0.000000	0.000000
25%	126.000000	1008.300000	34.000000	0.000000
50%	218.000000	1015.500000	66.600000	9.000000
75%	267.000000	1022.000000	88.900000	169.500000
max	360.000000	1048.900000	100.000000	1054.000000

	uvindex	solarenergy
count	94897.000000	94908.000000
mean	1.203104	0.439357
std	2.006343	0.715817
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	2.000000	0.600000
max	10.000000	3.800000

Counting for missing values

`print("Missing values:")`

`print(f_colchester_df.isnull().sum())`

Missing values:

datetime	0
temp	418
dew	424
humidity	422
precip	587
precipprob	434
preciptype	168576
snow	1546
snowdepth	1595
windspeed	397
winddir	429
sealevelpressure	7562
cloudcover	3619
solarradiation	93130
uvindex	93127
solarenergy	93116
dtype:	int64

Info about categorical variables

`print("Categories & frequencies for categorical variables:")`

```
for col in f_colchester_df.select_dtypes(include='object').columns:  
    print(f_colchester_df[col].value_counts())
```

Categories & frequencies for categorical variables:

```
2008-10-26T01:00:00    2  
2000-11-16T18:00:00    2  
2000-11-16T11:00:00    2  
2000-11-16T12:00:00    2  
2000-11-16T13:00:00    2
```

```
..  
2007-05-16T01:00:00    1  
2007-05-16T02:00:00    1  
2007-05-16T03:00:00    1  
2007-05-16T04:00:00    1  
2023-10-07T23:00:00    1
```

Name: datetime, Length: 187979, dtype: int64

```
rain    18543
```


```
rain,snow    851
```

```
snow    54
```

Name: preciptype, dtype: int64


```
# @title Precipitation
```

```
from matplotlib import pyplot as plt  
f_colchester_df['precip'].plot(kind='line', figsize=(10, 6),  
title='Precipitation')  
plt.gca().spines[['top', 'right']].set_visible(False)
```




```
# @title Temperature
```

```
from matplotlib import pyplot as plt
f_colchester_df['temp'].plot(kind='line', figsize=(8, 4),
title='temp')
plt.gca().spines[['top', 'right']].set_visible(False)
```



```
# @title Humidity
```

```
from matplotlib import pyplot as plt  
f_colchester_df['humidity'].plot(kind='line', figsize=(8, 4),  
title='humidity')  
plt.gca().spines[['top', 'right']].set_visible(False)
```




```
# @title DateTime VS Precipitation

from matplotlib import pyplot as plt
import seaborn as sns
def _plot_series(series, series_name, series_index=0):
    from matplotlib import pyplot as plt
    import seaborn as sns
    palette = list(sns.palettes.mpl_palette('rocket'))
    xs = series['datetime']
    ys = series['precip']

    plt.plot(xs, ys, label=series_name, color=palette[series_index %
len(palette)])

fig, ax = plt.subplots(figsize=(10, 5.2), layout='constrained')
df_sorted = f_colchester_df.sort_values('datetime', ascending=True)
for i, (series_name, series) in
enumerate(df_sorted.groupby('preciptype')):
    _plot_series(series, series_name, i)
    fig.legend(title='preciptype', bbox_to_anchor=(1, 1), loc='upper
left')
sns.despine(fig=fig, ax=ax)
plt.xlabel('Datetime')
_ = plt.ylabel('Precipitation')
```



```
# @title Time VS Dew

from matplotlib import pyplot as plt
import seaborn as sns
def _plot_series(series, series_name, series_index=0):
    from matplotlib import pyplot as plt
```


```

import seaborn as sns
palette = list(sns.palettes.mpl_palette('flare'))
xs = series['datetime']
ys = series['dew']

plt.plot(xs, ys, label=series_name, color=palette[series_index %
len(palette)])

fig, ax = plt.subplots(figsize=(10, 5.2), layout='constrained')
df_sorted = f_colchester_df.sort_values('datetime', ascending=True)
for i, (series_name, series) in
enumerate(df_sorted.groupby('preciptype')):
    _plot_series(series, series_name, i)
    fig.legend(title='preciptype', bbox_to_anchor=(1, 1), loc='upper
left')
sns.despine(fig=fig, ax=ax)
plt.xlabel('Time')
_ = plt.ylabel('Dew')

```



```

# @title Time VS Humidity


from matplotlib import pyplot as plt
import seaborn as sns
def _plot_series(series, series_name, series_index=0):
    from matplotlib import pyplot as plt
    import seaborn as sns
    palette = list(sns.palettes.mpl_palette('crest'))
    xs = series['datetime']
    ys = series['humidity']

    plt.plot(xs, ys, label=series_name, color=palette[series_index %

```

```
len(palette)])

fig, ax = plt.subplots(figsize=(10, 5.2), layout='constrained')
df_sorted = f_colchester_df.sort_values('datetime', ascending=True)
for i, (series_name, series) in
    enumerate(df_sorted.groupby('preciptype')):
        _plot_series(series, series_name, i)
        fig.legend(title='preciptype', bbox_to_anchor=(1, 1), loc='upper
left')
sns.despine(fig=fig, ax=ax)
plt.xlabel('Time')
_ = plt.ylabel('Humidity')
```



```
# @title Time vs Temperature

from matplotlib import pyplot as plt
import seaborn as sns
def _plot_series(series, series_name, series_index=0):
    from matplotlib import pyplot as plt
    import seaborn as sns
    palette = list(sns.palettes.mpl_palette('Dark2'))
    xs = series['datetime']
    ys = series['temp']


    plt.plot(xs, ys, label=series_name, color=palette[series_index %
len(palette)])

fig, ax = plt.subplots(figsize=(10, 5.2), layout='constrained')
df_sorted = f_colchester_df.sort_values('datetime', ascending=True)
for i, (series_name, series) in
    enumerate(df_sorted.groupby('preciptype')):
```

```

    _plot_series(series, series_name, i)
    fig.legend(title='precipdtype', bbox_to_anchor=(1, 1), loc='upper
left')
sns.despine(fig=fig, ax=ax)
plt.xlabel('Time')
_ = plt.ylabel('Temperature')

```



```


# @title Precipitaion VS Precipitation Probability

```


```

from matplotlib import pyplot as plt
sns.boxplot(x=f_colchester_df['precip'],
y=f_colchester_df['precipprob'])
<Axes: xlabel='precip', ylabel='precipprob'>

```





```
# @title Precipitaion VS Precipitation Probability
from matplotlib import pyplot as plt
f_colchester_df.plot(kind='scatter', x='precip', y='precipprob', s=32,
alpha=.8)
plt.gca().spines[['top', 'right']].set_visible(False)
```




```
# @title Temperature vs Dew
```

```
from matplotlib import pyplot as plt
f_colchester_df.plot(kind='box', x='temp')
plt.gca().spines[['top', 'right']].set_visible(False)
```




```
# @title Dew VS Humidity
```

```
from matplotlib import pyplot as plt
f_colchester_df.plot(kind='box', x='dew', y='humidity')
plt.gca().spines[['top', 'right']].set_visible(False)
```



```
# @title Humidity VS Precipitation


from matplotlib import pyplot as plt
f_colchester_df.plot(kind='scatter', x='humidity', y='precip', s=32,
alpha=.8)
plt.gca().spines[['top', 'right']].set_visible(False)
```



Exploring Important Columns


```
# @title Precipitation Type
```

```
from matplotlib import pyplot as plt
import seaborn as sns
f_colchester_df.groupby('preciptype').size().plot(kind='barh',
color=sns.palettes.mpl_palette('cubehelix'))
plt.gca().spines[['top', 'right']].set_visible(False)
```




```
# @title Temperature
```


```
from matplotlib import pyplot as plt
f_colchester_df['temp'].plot(kind='hist', bins=20, title='temp')
plt.gca().spines[['top', 'right']].set_visible(False)
```




```
# @title Dew  
from matplotlib import pyplot as plt  
f_colchester_df['dew'].plot(kind='hist', bins=20, title='dew')  
plt.gca().spines[['top', 'right',]].set_visible(False)
```



```
# @title Humidity  
from matplotlib import pyplot as plt  
f_colchester_df['humidity'].plot(kind='hist', bins=20,  
title='humidity')  
plt.gca().spines[['top', 'right',]].set_visible(False)
```

```
# @title Windspeed
from matplotlib import pyplot as plt
f_colchester_df['windspeed'].plot(kind='line', figsize=(8, 4),
title='windspeed')
plt.gca().spines[['top', 'right']].set_visible(False)
```



Time Series Analysis

```
months_of_the_year = ['January', 'February', 'March', 'April', 'May',
                       'June', 'July', 'August', 'September', 'October', 'November',
                       'December']
```

```
f_colchester_df.columns
```

```
Index(['datetime', 'temp', 'dew', 'humidity', 'precip', 'precipprob',
       'preciptype', 'snow', 'snowdepth', 'windspeed', 'winddir',
       'sealevelpressure', 'cloudcover', 'solarradiation', 'uvindex',
       'solarenergy'],
      dtype='object')
```

```
# Date will be our index. Let's convert it to a datetime type
```

```
f_colchester_df['datetime'] =
pd.to_datetime(f_colchester_df['datetime'], dayfirst=True)
f_colchester_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 188024 entries, 0 to 188023
Data columns (total 16 columns):
```

#	Column	Non-Null Count	Dtype
0	datetime	188024 non-null	datetime64[ns]
1	temp	187606 non-null	float64
2	dew	187600 non-null	float64
3	humidity	187602 non-null	float64
4	precip	187437 non-null	float64

```

5    precipprob      187590 non-null float64
6    preciptype      19448 non-null object
7    snow            186478 non-null float64
8    snowdepth       186429 non-null float64
9    windspeed       187627 non-null float64
10   winddir         187595 non-null float64
11   sealevelpressure 180462 non-null float64
12   cloudcover      184405 non-null float64
13   solarradiation   94894 non-null float64
14   uvindex         94897 non-null float64
15   solarenergy      94908 non-null float64
dtypes: datetime64[ns](1), float64(14), object(1)
memory usage: 23.0+ MB

```

```

df_c = f_colchester_df.copy()
df_c['month'] = df_c['datetime'].dt.month
df_c['year'] = df_c['datetime'].dt.year
df_c['week_of_year'] = df_c['datetime'].dt.isocalendar().week
df_c

```

```

{"type": "dataframe", "variable_name": "df_c"}

```

```

# Let's make the date column the index of the dataframe for easier
slicing
df_c.set_index('datetime', inplace=True) # note we can only run this
once, as it will delete the 'date' column.
df_c.head()

```

```

{"type": "dataframe", "variable_name": "df_c"}

```

```

print(df_c.isna().any())
df_c=df_c.dropna()
print(df_c.isna().any())

```

```


temp      True
dew        True
humidity   True
precip     True
precipprob True
preciptype True
snow       True
snowdepth  True
windspeed  True
winddir     True
sealevelpressure True
cloudcover  True
solarradiation True
uvindex     True
solarenergy True
month       False
year        False

```

```
week_of_year      False
dtype: bool
temp              False
dew               False
humidity          False
precip            False
precipprob        False
preciptype        False
snow              False
snowdepth         False
windspeed         False
winddir           False
sealevelpressure  False
cloudcover        False
solarradiation     False
uvindex           False
solarenergy        False
month             False
year              False
week_of_year      False
dtype: bool
```


```
plt.figure(figsize=(16,8))
plt.plot(df_c.index, df_c['temp'])
plt.xlabel('Year')
plt.ylabel('temperature (C)')
```

```
Text(0, 0.5, 'temperature (C)')
```



```
# Let's zoom in to 2014-2017
df_chunk = df_c.loc['2014-12':'2017-01'] # since the date is an
index, we can use it to filter our data

plt.figure(figsize=(10, 8))
plt.plot(df_chunk.index, df_chunk['temp'])
plt.xticks(rotation=90)
plt.xlabel('Date')
plt.ylabel('Average temperature (C)')
```



```
print(df_chunk.reindex(pd.date_range('2014-12', '2017-01')).isnull().all(1).sum())
df_chunk.reindex(pd.date_range('2014-12', '2017-01')).isnull().all(1)
```

652

```

2014-12-01    False
2014-12-02     True
2014-12-03    False
2014-12-04     True
2014-12-05    False
...
2016-12-28     True
2016-12-29     True
2016-12-30     True
2016-12-31     True
2017-01-01     True
Freq: D, Length: 763, dtype: bool

```

```
df_c[df_c.index.duplicated(keep=False)].head(20)
```

```

{"summary": "{\n  \"name\": \"df_c[df_c\", \n  \"rows\": 4, \n  \"fields\": [\n    {\n      \"column\": \"temp\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 1.2247448713915894, \n        \"min\": 11.9, \n        \"max\": 14.3, \n        \"samples\": [\n          14.0, \n          11.9, \n          14.3\n        ], \n        \"num_unique_values\": 4, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"dew\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 1.0144785195688797, \n        \"min\": 10.3, \n        \"max\": 12.6, \n        \"samples\": [\n          12.6, \n          10.3, \n          12.2\n        ], \n        \"num_unique_values\": 4, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"humidity\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 3.143897528016243, \n        \"min\": 87.17, \n        \"max\": 94.57, \n        \"samples\": [\n          91.71, \n          89.6, \n          87.17\n        ], \n        \"num_unique_values\": 4, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"precip\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 3.278237788812764, \n        \"min\": 0.197, \n        \"max\": 6.905, \n        \"samples\": [\n          0.197, \n          0.686, \n          0.21\n        ], \n        \"num_unique_values\": 4, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"precipprob\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.0, \n        \"min\": 100.0, \n        \"max\": 100.0, \n        \"samples\": [\n          100.0\n        ], \n        \"num_unique_values\": 1, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"preciptype\", \n      \"properties\": {\n        \"dtype\": \"category\", \n        \"samples\": [\n          \"rain\"\n        ], \n        \"num_unique_values\": 1, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"snow\", \n      \"properties\": {\n        \"dtype\": \"number\", \n

```

```

\"std\": 0.0,\n        \"min\": 0.0,\n        \"max\": 0.0,\n\"samples\": [\n        0.0\n    ],\n\"num_unique_values\": 1,\n    \"semantic_type\": \"\",\n\"description\": \"\"\n    },\n    {\n        \"column\":\n\"snowdepth\",\n        \"properties\": {\n            \"dtype\":\n\"number\",\n            \"std\": 0.0,\n            \"min\": 0.0,\n            \"max\": 0.0,\n            \"samples\": [\n                0.0\n            ],\n            \"num_unique_values\": 1,\n            \"semantic_type\": \"\",\n            \"description\": \"\"\n        },\n        {\n            \"column\":\n\"windspeed\",\n            \"properties\": {\n                \"dtype\":\n\"number\",\n                \"std\": 8.883083173463291,\n                \"min\": 15.0,\n                \"max\": 34.9,\n                \"samples\": [\n                    29.4\n                ],\n                \"num_unique_values\": 4,\n                \"semantic_type\": \"\",\n                \"description\": \"\"\n            },\n            {\n                \"column\": \"winddir\",\n                \"properties\": {\n                    \"dtype\": \"number\",\n                    \"std\": 13.32603967176045,\n                    \"min\": 192.0,\n                    \"max\": 223.0,\n                    \"samples\": [\n                        211.0\n                    ],\n                    \"num_unique_values\": 4,\n                    \"semantic_type\": \"\",\n                    \"description\": \"\"\n                },\n                {\n                    \"column\": \"sealevelpressure\",\n                    \"properties\": {\n                        \"dtype\":\n\"number\",\n                        \"std\": 9.184588540956335,\n                        \"min\": 981.5,\n                        \"max\": 997.9,\n                        \"samples\": [\n                            997.3\n                        ],\n                        \"num_unique_values\": 4,\n                        \"semantic_type\": \"\",\n                        \"description\": \"\"\n                    },\n                    {\n                        \"column\": \"cloudcover\",\n                        \"properties\": {\n                            \"dtype\": \"number\",\n                            \"std\": 9.40726669477732,\n                            \"min\": 71.9,\n                            \"max\": 93.8,\n                            \"samples\": [\n                                84.4\n                            ],\n                            \"num_unique_values\": 4,\n                            \"semantic_type\": \"\",\n                            \"description\": \"\"\n                        },\n                        {\n                            \"column\": \"solarradiation\",\n                            \"properties\": {\n                                \"dtype\":\n\"number\",\n                                \"std\": 0.0,\n                                \"min\": 0.0,\n                                \"max\": 0.0,\n                                \"samples\": [\n                                    0.0\n                                ],\n                                \"num_unique_values\": 1,\n                                \"semantic_type\": \"\",\n                                \"description\": \"\"\n                            },\n                            {\n                                \"column\": \"uvindex\",\n                                \"properties\": {\n                                    \"dtype\": \"number\",\n                                    \"std\": 0.0,\n                                    \"min\": 0.0,\n                                    \"max\": 0.0,\n                                    \"samples\": [\n                                        0.0\n                                    ],\n                                    \"num_unique_values\": 1,\n                                    \"semantic_type\": \"\",\n                                    \"description\": \"\"\n                                },\n                                {\n                                    \"column\": \"solarenergy\",\n                                    \"properties\": {\n                                        \"dtype\":\n\"number\",\n                                        \"std\": 0.0,\n                                        \"min\": 0.0,\n                                        \"max\": 0.0,\n                                        \"samples\": [\n                                            0.0\n                                        ],\n                                        \"num_unique_values\": 1,\n                                        \"semantic_type\": \"\",\n                                        \"description\": \"\"\n                                    },\n                                    {\n                                        \"column\": \"month\",\n                                        \"properties\": {\n                                            \"dtype\": \"number\",\n                                            \"std\": 0,\n                                            \"min\": 10,\n                                            \"max\": 10,\n                                            \"samples\": [\n                                                10\n                                            ],\n
```

```

\"num_unique_values\": 1,\n      \"semantic_type\": \"\",\n\"description\": \"\"\n    },\n    {\n      \"column\":\n\"year\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 5,\n        \"min\": 2013,\n        \"max\": 2023,\n        \"samples\": [\n          2023\n        ],\n        \"num_unique_values\": 2,\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      {\n        \"column\":\n\"week_of_year\",\n        \"properties\": {\n          \"dtype\":\n\"UInt32\",\n          \"samples\": [\n            43\n          ],\n          \"num_unique_values\": 1,\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        }\n      }\n    ],\n  },\n  \"type\": \"dataframe\"}

```

Let's keep the first one only - in practice this would require more careful analysis!

```

df_c = df_c[~df_c.index.duplicated(keep='first')]
len(df_c)

```

11497

Now we can reindex -- this is where the original error about duplicates was

```

df_c = df_c.reindex(pd.date_range(df_c.index[0], df_c.index[-1]))
print(len(df_c))

```

4996

Now we should have missing values

```

print(df_c.isna().sum())

```


```

temp          4339
dew            4339
humidity       4339
precip         4339
precipprob     4339
preciptype     4339
snow           4339
snowdepth      4339
windspeed      4339
winddir        4339
sealevelpressure 4339
cloudcover     4339
solarradiation 4339
uvindex        4339
solarenergy    4339
month          4339
year           4339
week_of_year   4339
dtype: int64

```

df_chunk = df_c.loc['2017-12-15':'2018-01-15'] # since the date is an index, we can use it to filter our data



```
plt.figure(figsize=(10, 8))
plt.plot(df_chunk.index, df_chunk['temp'])
plt.xticks(rotation=90)
plt.xlabel('Date')
_=plt.ylabel('temperature (C)')
# The missing values are clearly visible now!
```



```
df2 = df_chunk.copy()
df2 = df2.loc[:, 'temp'].to_frame()
df2

{"summary": "{\n  \"name\": \"df2\",\n  \"rows\": 32,\n  \"fields\": [\n    {\n      \"column\": \"temp\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 2.311038172747788,\n        \"min\": 1.8,\n        \"max\": 11.5,\n        \"samples\": [\n
```

[illegible]



Try different ways to fill the data - more advanced: interpolation

```
df2['linear_interp'] = df2['temp'].interpolate(method='linear')
df2['nearest_interp'] = df2['temp'].interpolate(method='nearest')
df2['spline_interp'] = df2['temp'].interpolate(method='spline',
order=2)
df2['polynomial_interp'] =
df2['temp'].interpolate(method="polynomial", order=3)
```

Plot

```
fig, ax = plt.subplots(figsize=(10,8))


plt.plot(df2.index, df2['linear_interp'], linestyle='--', color='red',
label='linear')
plt.plot(df2.index, df2['nearest_interp'], linestyle='-.',
color='green', label='nearest')
plt.plot(df2.index, df2['spline_interp'], linestyle=':',
color='black', label='spline')
plt.plot(df2.index, df2['polynomial_interp'], linestyle='--',
```

```

color='purple', label='polynomial')
plt.plot(df2.index, df2['temp'], label='Original', color='blue')

plt.legend()
plt.ylabel('Temperature')
plt.ylim(-1, 26)
_=plt.title('Interpolation')


```




```

plt.figure(figsize=(16, 8))
_=sns.boxplot(x='year', y='temp', data=df_c)
_=plt.tight_layout()


```



```
# Visualise trends across years
sns.lineplot(x='month', y='temp', data=df, hue='year')
_ = plt.xticks(np.arange(1, 13), months_of_the_year, rotation=90)
_ = plt.xlim(1, 12) # limit x-axis
_ = plt.tight_layout()
```




```
# Visualise trends across years
fix, ax = plt.subplots(2, 1, sharex=True, figsize=(20,15))
sns.boxplot(x='month', y='temp', data=df, ax=ax[0]) # top plot
sns.boxplot(x='month', y='temp', data=df, hue='year', ax=ax[1]) #
bottom plot
ax[1].set_xticks(np.arange(0, 12), months_of_the_year, fontsize=14)
plt.tight_layout()
```



```
data_ds = df['temp'].resample('M').mean().ffill().to_frame() # one
value per month
data_ds

{"summary":{"\n  \"name\": \"data_ds\", \n  \"rows\": 169, \n  \"fields\": [\n    {\n      \"column\": \"temp\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 3.770517588760982, \n        \"min\": 2.218867924528302, \n        \"max\": 18.932758620689654, \n        \"samples\": [\n          14.439473684210528, \n          16.96818181818182, \n          4.944067796610169\n        ], \n        \"num_unique_values\": 125, \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    }\n  ], \n  \"type\": \"dataframe\", \n  \"variable_name\": \"data_ds\"}

# Try decomposition on the resampled dataset
from statsmodels.tsa.seasonal import seasonal_decompose, STL
decomposition = STL(data_ds['temp']).fit()
fig = decomposition.plot()
fig.set_size_inches(12,8)
fig.axes[1].grid()
```



```
# Statistical test for stationarity: Augmented Dickey-Fuller (ADF) test
```

```
adf_result = adfuller(data_ds['temp'])
print('ADF Statistic %.2f:' % adf_result[0])
print('ADF p-value: %.4f:' % adf_result[1])
# p-value < 0.05 ==> timeseries does not have a unit root
```

```
ADF Statistic -2.17:
ADF p-value: 0.2180:
```


```
adf_result = adfuller(data_ds.loc['2017':'2024', 'temp']) # ADF test on the full years only. Is there a trend?
```

```
print('ADF Statistic %.2f:' % adf_result[0])
print('ADF p-value: %.4f:' % adf_result[1])
```


```
ADF Statistic -1.51:
ADF p-value: 0.5292:
```

```
# Autocorrelation (can help us with modelling later)
```

```
fig, ax = plt.subplots(figsize=(16,8))
_ = plot_acf(data_ds['temp'], lags=48, ax=ax) # each lag is one month, so we're looking at 4 years worth of past data
_ = plt.xlabel('Lags (months)')
_ = plt.ylabel('Autocorrelation')
```

```
# Partial autocorrelation (can help us with modelling later)
fig, ax = plt.subplots(figsize=(16,8))
_ = plot_pacf(data_ds['temp'], ax=ax)
_ = plt.xlabel('Lags (months)')
_ = plt.ylabel('Partial autocorrelation')
```



```
# https://seaborn.pydata.org/examples/many\_pairwise\_correlations.html
# Compute the correlation matrix
```

```

corr = df.iloc[:, :-3].corr()

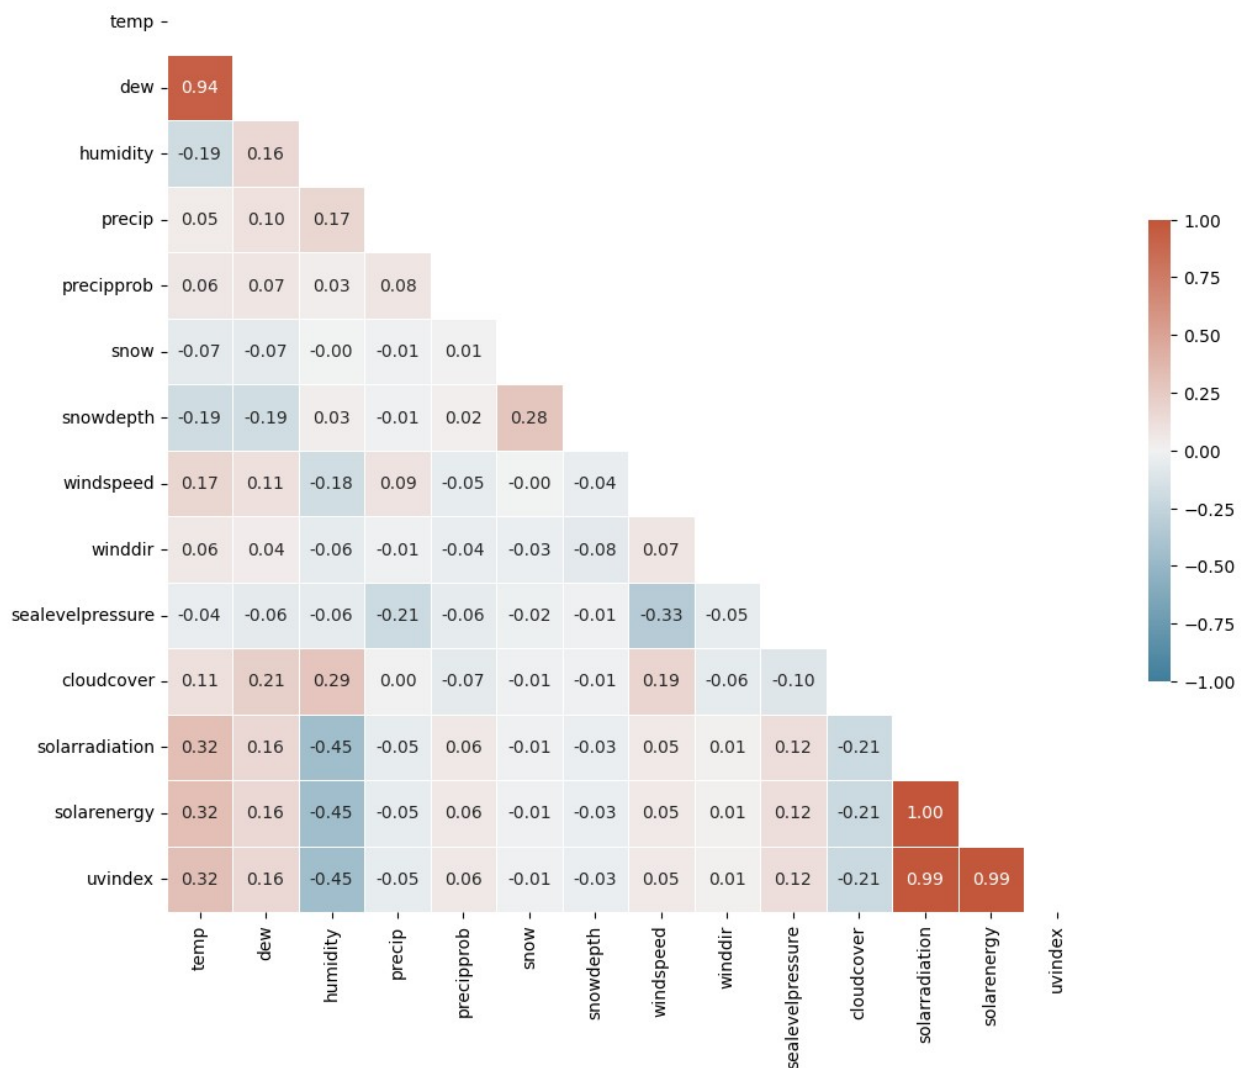
# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr, dtype=bool))

# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(11, 9))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(230, 20, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmin=-1, vmax=1, center=0,
            annot=True, square=True, linewidths=.5, cbar_kws={"shrink": .5},
            fmt='.2f')
plt.tight_layout()

```



```

# Convert columns to numeric type if necessary
df_numeric = df.iloc[:, :-3].apply(pd.to_numeric, errors='coerce')

# Compute the correlation matrix
corr_diff = df_numeric.diff().corr()
# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr_diff, dtype=bool))

# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(11, 9))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(230, 20, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr_diff, mask=mask, cmap=cmap, vmin=-1, vmax=1,
            center=0, annot=True,
            square=True, linewidths=.5, cbar_kws={"shrink": .5},
            fmt='.2f')
plt.tight_layout()
plt.show()

```



```

{\n      \"dtype\": \"number\", \n      \"std\": 8.888867917418999, \n      \"min\": 41.89, \n      \"max\": 100.0, \n      \"samples\": [\n        72.53, \n        93.24, \n        98.62 \n      ], \n      \"num_unique_values\": 2802, \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    }, \n    {\n      \"column\": \"precip\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 1.894431542189717, \n        \"min\": 0.0, \n        \"max\": 32.385, \n        \"samples\": [\n          0.262, \n          0.074, \n          0.56 \n        ], \n        \"num_unique_values\": 2216, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"precipprob\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 17.376592123895172, \n        \"min\": 0.0, \n        \"max\": 100.0, \n        \"samples\": [\n          0.0, \n          100.0 \n        ], \n        \"num_unique_values\": 2, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"preciptype\", \n      \"properties\": {\n        \"dtype\": \"category\", \n        \"samples\": [\n          \"rain\", \n          \"rain,snow\" \n        ], \n        \"num_unique_values\": 3, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"snow\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.0025808656410961357, \n        \"min\": 0.0, \n        \"max\": 0.13, \n        \"samples\": [\n          0.08, \n          0.13 \n        ], \n        \"num_unique_values\": 5, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"snowdepth\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.3505657600610243, \n        \"min\": 0.0, \n        \"max\": 14.0, \n        \"samples\": [\n          4.91, \n          0.2 \n        ], \n        \"num_unique_values\": 39, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"windspeed\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 10.354676031208625, \n        \"min\": 0.4, \n        \"max\": 72.2, \n        \"samples\": [\n          12.0, \n          39.0 \n        ], \n        \"num_unique_values\": 501, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"winddir\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 89.50662453559492, \n        \"min\": 1.0, \n        \"max\": 360.0, \n        \"samples\": [\n          136.0, \n          15.0 \n        ], \n        \"num_unique_values\": 361, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"sealevelpressure\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 11.43377469655269, \n        \"min\": 955.0, \n        \"max\": 1048.5, \n        \"samples\": [\n          990.2, \n          1009.2 \n        ], \n        \"num_unique_values\": 669, \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      \"column\": \"cloudcover\", \n

```

```

{"properties": {"dtype": "number", "std": 26.483150036400094, "min": 0.0, "max": 100.0, "samples": [36.4, 10.6], "num_unique_values": 913, "semantic_type": "", "description": ""}, {"column": "solarradiation", "properties": {"dtype": "number", "std": 158.6456589417339, "min": 0.0, "max": 917.0, "samples": [248.1, 365.0], "num_unique_values": 1720, "semantic_type": "", "description": ""}, {"column": "solarenergy", "properties": {"dtype": "number", "std": 0.571534271107572, "min": 0.0, "max": 3.3, "samples": [0.8, 2.3], "num_unique_values": 34, "semantic_type": "", "description": ""}, {"column": "uvindex", "properties": {"dtype": "number", "std": 1.6080207673456206, "min": 0.0, "max": 9.0, "samples": [6.0, 2.0], "num_unique_values": 10, "semantic_type": "", "description": ""}, {"column": "temp_avg_lag3", "properties": {"dtype": "number", "std": 4.639795491228626, "min": -8.9, "max": 28.3, "samples": [8.2, 13.4], "num_unique_values": 281, "semantic_type": "", "description": ""}}], "type": "dataframe", "variable_name": "df2"}

```

Let's see what happens if we do the differential operation again.

```

df2 = df.iloc[:, :-3].copy()
df2['temp_avg_lag3'] = df2['temp'].shift(-3)
corr2 = df2.corr()
# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr2, dtype=bool))

```

Set up the matplotlib figure

```
f, ax = plt.subplots(figsize=(11, 9))
```

Generate a custom diverging colormap

```
cmap = sns.diverging_palette(230, 20, as_cmap=True)
```

Draw the heatmap with the mask and correct aspect ratio

```

sns.heatmap(corr2, mask=mask, cmap=cmap, vmin=-1, vmax=1, center=0,
            annot=True, square=True, linewidths=.5, cbar_kws={"shrink": .5},
            fmt='.2f')
plt.tight_layout()

```

